

Exploratory Data Analysis Using Python and BigQuery

Learning Objectives

1. Analyze a Pandas Dataframe
2. Create Seaborn plots for Exploratory Data Analysis in Python
3. Write a SQL query to pick up specific fields from a BigQuery dataset
4. Exploratory Analysis in BigQuery

Introduction

This lab is an introduction to linear regression using Python and Scikit-Learn. This lab serves as a foundation for more complex algorithms and machine learning models that you will encounter in the course. We will train a linear regression model to predict housing price.

Each learning objective will correspond to a **#TODO** in the [student lab notebook](#) -- try to complete that notebook first before reviewing this solution notebook.

Import Libraries

```
In [ ]: # Run the chown command to change the ownership
!sudo chown -R jupyter:jupyter /home/jupyter/training-data-analyst
```

```
In [ ]: # Install the Google Cloud BigQuery Library
!pip install --user google-cloud-bigquery==1.25.0
```

Please ignore any incompatibility warnings and errors.

Restart the kernel before proceeding further (On the Notebook menu - Kernel - Restart Kernel).

```
In [2]: # You can use any Python source file as a module by executing an import statement in some other Python source file.
# The import statement combines two operations; it searches for the named module, then it binds the results of that search
# to a name in the local scope.
import os
import pandas as pd
import numpy as np
# Import matplotlib to visualize the model
import matplotlib.pyplot as plt
# Seaborn is a Python data visualization library based on matplotlib
import seaborn as sns
%matplotlib inline
```

Load the Dataset

Here, we create a directory called usahousing. This directory will hold the dataset that we copy from Google Cloud Storage.

```
In [3]: # Create a directory to hold the dataset
if not os.path.isdir("../data/explore"):
    os.makedirs("../data/explore")
```

Next, we copy the Usahousing dataset from Google Cloud Storage.

```
In [4]: # Copy the file using `gsutil cp` from Google Cloud Storage in the required directory
!gsutil cp gs://cloud-training-demos/feat_eng/housing/housing_pre-proc.csv ../data/explore
```

```
Copying gs://cloud-training-demos/feat_eng/housing/housing_pre-proc.csv...
/ [1 files] 1.4 MiB / 1.4 MiB
Operation completed over 1 objects/1.4 MiB.
```

Then we use the "ls" command to list files in the directory. This ensures that the dataset was copied.

```
In [5]: # `ls` shows the working directory's contents.
# The `l` flag list the all files with permissions and details
!ls -l ../data/explore
```

```
total 1404
-rw-r--r-- 1 jupyter jupyter 1435069 Jun 24 21:24 housing_pre-proc.csv
```

Next, we read the dataset into a Pandas dataframe.

```
In [6]: # TODO 1
# Read a comma-separated values (csv) file into a DataFrame using the read_csv() function
```

```
df_USAhousing = pd.read_csv('../data/explore/housing_pre-proc.csv')
```

Inspect the Data

```
In [7]: # Get the first five rows using the head() method

df_USAhousing.head()
```

```
Out[7]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0	NEAR BAY
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0	NEAR BAY
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0	NEAR BAY

Let's check for any null values.

```
In [8]: # `isnull()` finds a null value in a column and `sum()` counts it
df_USAhousing.isnull().sum()
```

```
Out[8]: longitude      0
latitude      0
housing_median_age  0
total_rooms    0
total_bedrooms  0
population     0
households     0
median_income  0
median_house_value  0
ocean_proximity  0
dtype: int64
```

```
In [9]: # Get some basic statistical details using describe() method
df_stats = df_USAhousing.describe()
# Transpose index and columns of the dataframe
df_stats = df_stats.transpose()
df_stats
```

```
Out[9]:
```

	count	mean	std	min	25%	50%	75%	max
longitude	20433.0	-119.570689	2.003578	-124.3500	-121.8000	-118.4900	-118.010	-114.3100
latitude	20433.0	35.633221	2.136348	32.5400	33.9300	34.2600	37.720	41.9500
housing_median_age	20433.0	28.633094	12.591805	1.0000	18.0000	29.0000	37.000	52.0000
total_rooms	20433.0	2636.504233	2185.269567	2.0000	1450.0000	2127.0000	3143.000	39320.0000
total_bedrooms	20433.0	537.870553	421.385070	1.0000	296.0000	435.0000	647.000	6445.0000
population	20433.0	1424.946949	1133.208490	3.0000	787.0000	1166.0000	1722.000	35682.0000
households	20433.0	499.433465	382.299226	1.0000	280.0000	409.0000	604.000	6082.0000
median_income	20433.0	3.871162	1.899291	0.4999	2.5637	3.5365	4.744	15.0001
median_house_value	20433.0	206864.413155	115435.667099	14999.0000	119500.0000	179700.0000	264700.000	500001.0000

```
In [10]: # Get a concise summary of a DataFrame
df_USAhousing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20433 entries, 0 to 20432
Data columns (total 10 columns):
longitude      20433 non-null float64
latitude      20433 non-null float64
housing_median_age  20433 non-null float64
total_rooms    20433 non-null float64
total_bedrooms 20433 non-null float64
population     20433 non-null float64
households     20433 non-null float64
median_income  20433 non-null float64
median_house_value 20433 non-null float64
ocean_proximity 20433 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

Let's take a peek at the first and last five rows of the data for all columns.

```
In [11]:
```

```
print ("Rows      : ",df_USAhousing.shape[0])
print ("Columns   : ",df_USAhousing.shape[1])
print ("\nFeatures : \n",df_USAhousing.columns.tolist())
print ("\nMissing values : ", df_USAhousing.isnull().sum().values.sum())
print ("\nUnique values : \n",df_USAhousing
      .nunique())
```

```
Rows      : 20433
Columns   : 10
```

```
Features :
['longitude', 'latitude', 'housing_median_age', 'total_rooms', 'total_bedrooms', 'population', 'households', 'median_income', 'median_house_value', 'ocean_proximity']
```

```
Missing values : 0
```

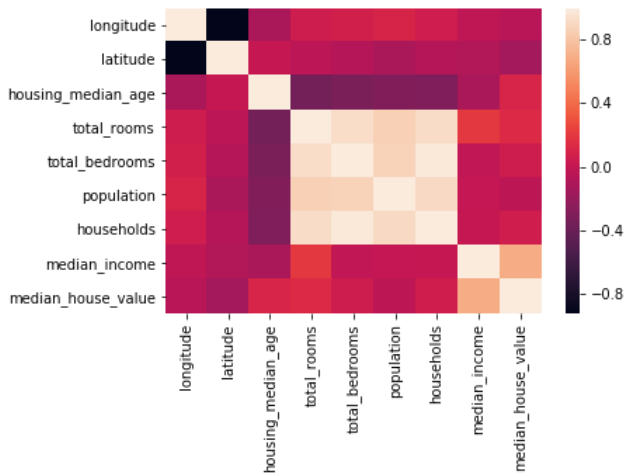
```
Unique values :
longitude      844
latitude       861
housing_median_age  52
total_rooms    5911
total_bedrooms 1923
population     3879
households     1809
median_income  12825
median_house_value 3833
ocean_proximity 5
dtype: int64
```

Explore the Data

Let's create some simple plots to check out the data!

```
In [12]: # `heatmap` plots a rectangular data in a color-encoded matrix and
# `corr` finds the pairwise correlation of all columns in the dataframe
sns.heatmap(df_USAhousing.corr())
```

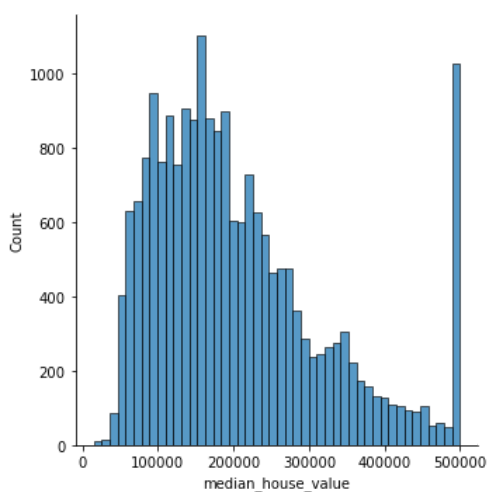
```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7f25ba7c1dd8>
```



Create a distplot showing "median_house_value".

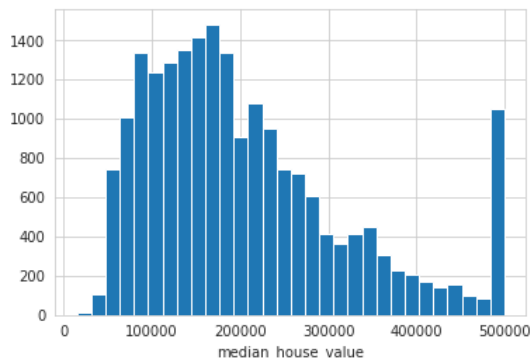
```
In [13]: # TODO 2a
# Plot a univariate distribution of observations using seaborn `distplot()` function
sns.distplot(df_USAhousing['median_house_value'])
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f25ba6d15f8>
```



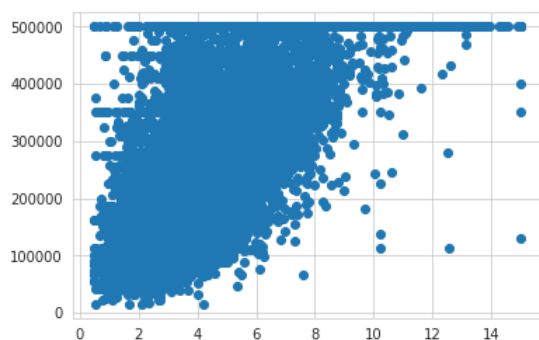
```
In [14]: # Set the aesthetic style of the plots
sns.set_style('whitegrid')
# Plot a histogram using `hist()` function
df_USAhousing['median_house_value'].hist(bins=30)
plt.xlabel('median_house_value')
```

Out[14]: Text(0.5, 0, 'median_house_value')



```
In [15]: x = df_USAhousing['median_income']
y = df_USAhousing['median_house_value']

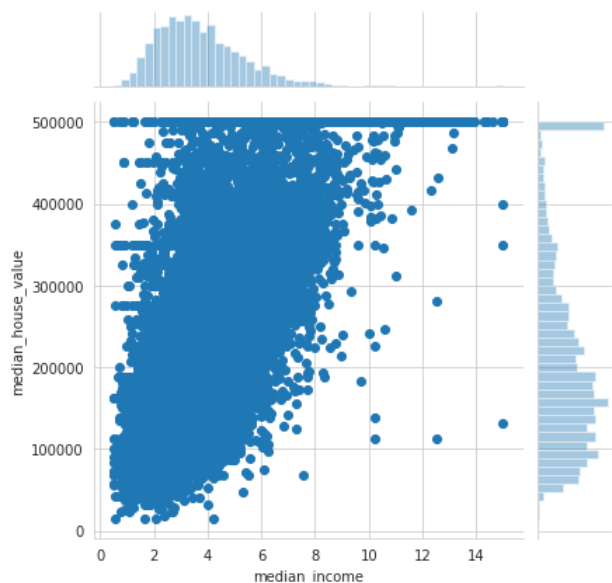
# Scatter plot of y vs x using scatter() and `show()` display all open figures
plt.scatter(x, y)
plt.show()
```



Create a jointplot showing "median_income" versus "median_house_value".

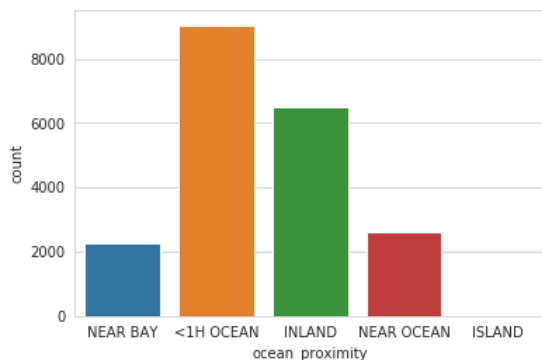
```
In [16]: # TODO 2b
# `jointplot()` draws a plot of two variables with bivariate and univariate graphs.
sns.jointplot(x='median_income', y='median_house_value', data=df_USAhousing)
```

Out[16]: <seaborn.axisgrid.JointGrid at 0x7f25ba74a630>

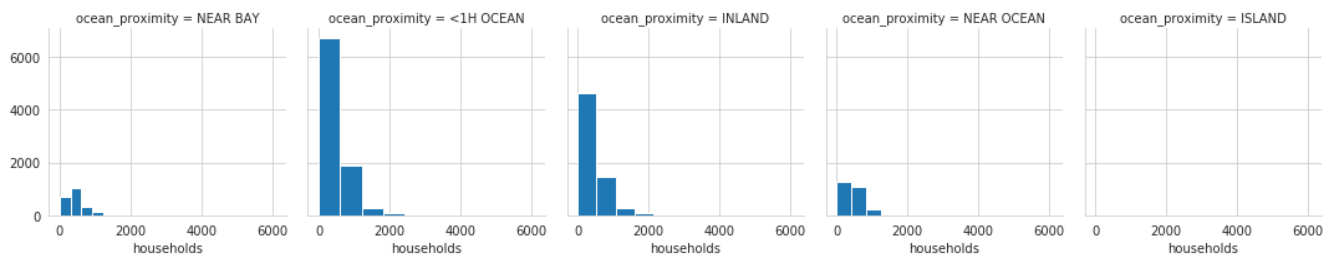


```
In [17]: # `countplot()` shows the counts of observations in each categorical bin using bars
sns.countplot(x = 'ocean_proximity', data=df_USAhousing)
```

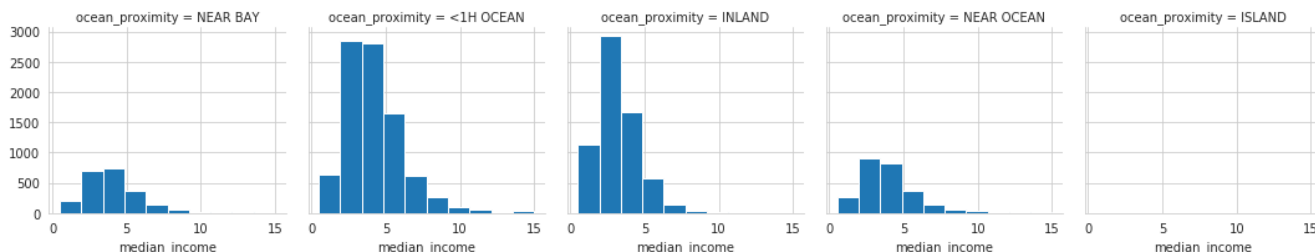
```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x7f25ba4ef400>
```



```
In [18]: # takes numeric only?
# plt.figure(figsize=(20,20))
# Draw a multi-plot on every facet using `FacetGrid()`
g = sns.FacetGrid(df_USAhousing, col="ocean_proximity")
# Pass a function and the name of one or more columns in the dataframe
g.map(plt.hist, "households");
```



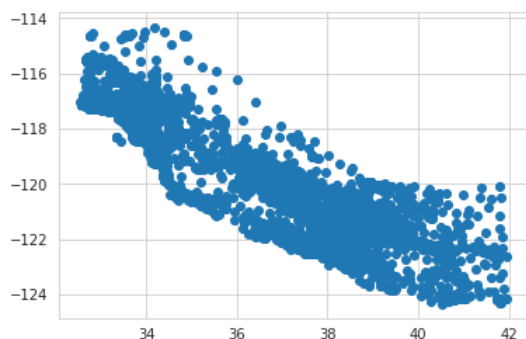
```
In [60]: # takes numeric only?
# plt.figure(figsize=(20,20))
# Draw a multi-plot on every facet using `FacetGrid()`
g = sns.FacetGrid(df_USAhousing, col="ocean_proximity")
# Pass a function and the name of one or more columns in the dataframe
g.map(plt.hist, "median_income");
```



You can see below that this is the state of California!

```
In [20]: x = df_USAhousing['latitude']
y = df_USAhousing['longitude']

# Scatter plot of y vs x and display all open figures
plt.scatter(x, y)
plt.show()
```



Explore and create ML datasets

In this notebook, we will explore data corresponding to taxi rides in New York City to build a Machine Learning model in support of a fare-estimation tool. The idea is to suggest a likely fare to taxi riders so that they are not surprised, and so that they can protest if the charge is much higher than expected.

Learning Objectives

- Access and explore a public BigQuery dataset on NYC Taxi Cab rides
- Visualize your dataset using the Seaborn library

First, **restart the Kernel**. Now, let's start with the Python imports that we need.

```
In [1]: # Import the python Libraries
from google.cloud import bigquery
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
```

Extract sample data from BigQuery

The dataset that we will use is [a BigQuery public dataset](#). Click on the link, and look at the column names. Switch to the Details tab to verify that the number of records is one billion, and then switch to the Preview tab to look at a few rows.

Let's write a SQL query to pick up interesting fields from the dataset. It's a good idea to get the timestamp in a predictable format.

```
In [2]: %%bigquery
# SQL query to get a fields from dataset which prints the 10 records
SELECT
  FORMAT_TIMESTAMP(
    "%Y-%m-%d %H:%M:%S %Z", pickup_datetime) AS pickup_datetime,
  pickup_longitude, pickup_latitude, dropoff_longitude,
  dropoff_latitude, passenger_count, trip_distance, tolls_amount,
  fare_amount, total_amount
# TODO 3
FROM
  `nyc-tlc.yellow.trips`
LIMIT 10
```

Out[2]:

	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	trip_distance	tolls_amount	fare_amount	total
0	2010-03-23 16:04:00 UTC	-73.937463	40.758173	-73.937462	40.758175	1	0.00	0.0	0.0	
1	2010-03-31 19:34:22 UTC	-73.848701	40.738563	-73.980793	40.764666	1	83.60	0.0	0.0	
2	2010-02-18 23:22:37 UTC	-74.045752	40.720509	-74.045752	40.720509	3	0.00	0.0	0.0	
3	2015-01-10 15:20:08 UTC	-73.972122	40.760014	-73.971947	40.760063	1	0.05	0.0	0.0	
4	2010-03-26 21:44:51 UTC	-73.989504	40.730073	-73.915033	40.766109	1	6.00	0.0	0.0	
5	2010-02-26 21:15:13 UTC	-73.790439	40.644783	-74.036504	40.722646	2	218.00	0.0	0.0	
6	2010-02-14 23:41:37 UTC	-73.862951	40.734516	-73.982297	40.761860	1	80.00	0.0	0.0	
7	2010-03-27 14:46:11 UTC	-73.980504	40.748292	-73.980181	40.748143	1	0.00	0.0	0.0	
8	2010-02-04 01:00:03 UTC	-73.903466	40.747651	-73.903396	40.747749	1	0.30	0.0	0.0	
9	2013-08-01 01:23:36 UTC	-73.980993	40.764660	-73.981463	40.764557	1	2.80	0.0	0.0	

Let's increase the number of records so that we can do some neat graphs. There is no guarantee about the order in which records are returned, and so no guarantee about which records get returned if we simply increase the LIMIT. To properly sample the dataset, let's use the HASH of the pickup time and return 1 in 100,000 records -- because there are 1 billion records in the data, we should get back approximately 10,000 records if we do this.

We will also store the BigQuery result in a Pandas dataframe named "trips"

In [3]:

```
%%bigquery trips
# SQL query to save the results in the trips dataframe
SELECT
  FORMAT_TIMESTAMP(
    "%Y-%m-%d %H:%M:%S %Z", pickup_datetime) AS pickup_datetime,
  pickup_longitude, pickup_latitude,
  dropoff_longitude, dropoff_latitude,
  passenger_count,
  trip_distance,
  tolls_amount,
  fare_amount,
  total_amount
FROM
  `nyc-tlc.yellow.trips`
WHERE
  ABS(MOD(FARM_FINGERPRINT(CAST(pickup_datetime AS STRING)), 100000)) = 1
```

Out[3]:

	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	trip_distance	tolls_amount	fare_amount	
0	2011-06-04 02:52:10 UTC	-73.984681	40.769893	-74.007312	40.705326	1	5.30	0.00	15.3	
1	2014-10-06 15:16:00 UTC	-73.980130	40.760910	-73.861730	40.768330	2	11.47	5.33	36.5	
2	2014-12-08 21:50:00 UTC	-73.870867	40.773782	-74.003297	40.708215	2	11.81	0.00	33.5	
3	2010-05-26 16:15:03 UTC	-74.002922	40.714474	-73.978505	40.758280	1	6.10	0.00	20.9	
4	2012-05-05 22:46:05 UTC	-74.009790	40.712483	-73.959293	40.768908	1	5.20	0.00	16.9	
...
10784	2010-02-25 20:14:00 UTC	-73.991358	40.749645	-73.971768	40.792898	5	3.84	0.00	12.9	
10785	2012-06-06 12:34:32 UTC	-73.946107	40.778459	-73.975723	40.760311	1	2.80	0.00	12.9	
10786	2010-05-06 18:16:06 UTC	-73.989262	40.741876	-73.951123	40.782685	1	4.00	0.00	12.9	
10787	2012-08-10 12:17:02 UTC	-74.015546	40.707420	-73.990367	40.729593	1	2.90	0.00	12.9	

	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	trip_distance	tolls_amount	fare_amount
10788	2011-08-11 23:35:42 UTC	-73.978100	40.725390	-73.992485	40.698344	1	4.00	0.00	12.9

10789 rows × 10 columns



In [4]: `print(len(trips))`

10789

In [5]: `# We can slice Pandas dataframes as if they were arrays`
`trips[:10]`

Out[5]:

	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	trip_distance	tolls_amount	fare_amount	total
0	2011-06-04 02:52:10 UTC	-73.984681	40.769893	-74.007312	40.705326	1	5.30	0.00	15.3	
1	2014-10-06 15:16:00 UTC	-73.980130	40.760910	-73.861730	40.768330	2	11.47	5.33	36.5	
2	2014-12-08 21:50:00 UTC	-73.870867	40.773782	-74.003297	40.708215	2	11.81	0.00	33.5	
3	2010-05-26 16:15:03 UTC	-74.002922	40.714474	-73.978505	40.758280	1	6.10	0.00	20.9	
4	2012-05-05 22:46:05 UTC	-74.009790	40.712483	-73.959293	40.768908	1	5.20	0.00	16.9	
5	2010-12-21 13:08:00 UTC	-73.982422	40.739847	-73.981658	40.768732	2	2.64	0.00	14.9	
6	2011-12-03 10:28:00 UTC	-73.998822	40.680933	-73.968960	40.757878	1	8.28	0.00	20.9	
7	2014-05-20 23:09:00 UTC	-73.995203	40.727307	-73.948775	40.813487	1	10.31	0.00	33.5	
8	2010-10-09 23:50:28 UTC	-73.956644	40.771152	-74.005279	40.740280	1	10.10	0.00	25.7	
9	2012-07-05 14:18:00 UTC	-73.916517	40.743212	-73.956785	40.780882	1	4.74	0.00	15.7	

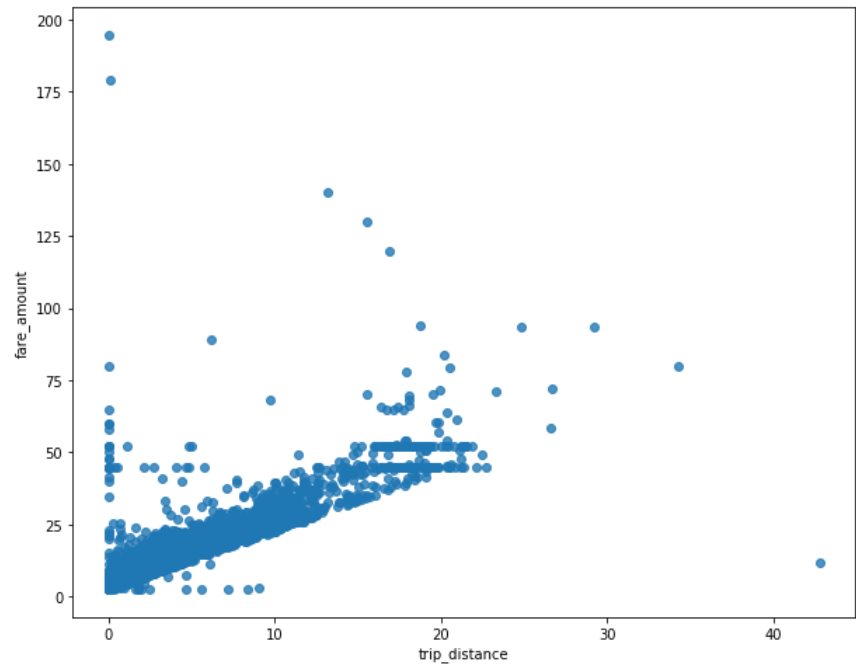


Exploring data

Let's explore this dataset and clean it up as necessary. We'll use the Python Seaborn package to visualize graphs and Pandas to do the slicing and filtering.

In [6]:

```
# TODO 4
# Use Seaborn `regplot()` function to plot the data and a linear regression model fit.
ax = sns.regplot(
    x="trip_distance", y="fare_amount",
    fit_reg=False, ci=None, truncate=True, data=trips)
ax.figure.set_size_inches(10, 8)
```

Hmm ... do you see something wrong with the data that needs addressing?

It appears that we have a lot of invalid data that is being coded as zero distance and some fare amounts that are definitely illegitimate. Let's remove them from our analysis. We can do this by modifying the BigQuery query to keep only trips longer than zero miles and fare amounts that are at least the minimum cab fare (\$2.50).
Note the extra WHERE clauses.

```
In [7]: %%bigquery trips
# SQL query with where clause to save the results in the trips dataframe
SELECT
  FORMAT_TIMESTAMP(
    "%Y-%m-%d %H:%M:%S %Z", pickup_datetime) AS pickup_datetime,
  pickup_longitude, pickup_latitude,
  dropoff_longitude, dropoff_latitude,
  passenger_count,
  trip_distance,
  tolls_amount,
  fare_amount,
  total_amount
FROM
  `nyc-tlc.yellow.trips`
WHERE
  ABS(MOD(FARM_FINGERPRINT(CAST(pickup_datetime AS STRING)), 100000)) = 1
# TODO 4a
  AND trip_distance > 0
  AND fare_amount >= 2.5
```

	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	trip_distance	tolls_amount	fare_amount
0	2014-10-06 15:16:00 UTC	-73.980130	40.760910	-73.861730	40.768330	2	11.47	5.33	36.5
1	2014-12-08 21:50:00 UTC	-73.870867	40.773782	-74.003297	40.708215	2	11.81	0.00	33.5
2	2010-05-26 16:15:03 UTC	-74.002922	40.714474	-73.978505	40.758280	1	6.10	0.00	20.9
3	2012-05-05 22:46:05 UTC	-74.009790	40.712483	-73.959293	40.768908	1	5.20	0.00	16.9
4	2010-12-21 13:08:00 UTC	-73.982422	40.739847	-73.981658	40.768732	2	2.64	0.00	14.9
...
10711	2010-02-25 20:14:00 UTC	-73.991358	40.749645	-73.971768	40.792898	5	3.84	0.00	12.9
10712	2012-06-06 12:34:32 UTC	-73.946107	40.778459	-73.975723	40.760311	1	2.80	0.00	12.9
10713	2011-10-27 07:55:01 UTC	-73.991334	40.749870	-73.954420	40.764275	1	3.20	0.00	12.9

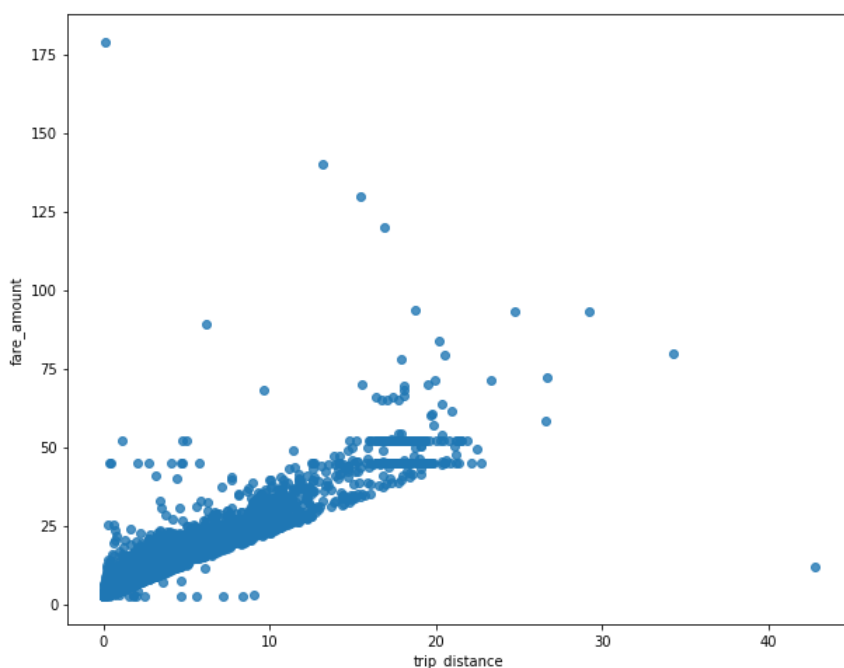
	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	trip_distance	tolls_amount	fare_amount
10714	2010-05-06 18:16:06 UTC	-73.989262	40.741876	-73.951123	40.782685	1	4.00	0.00	12.9
10715	2009-09-25 03:47:00 UTC	-73.991533	40.735107	-73.964660	40.687678	1	4.69	0.00	12.9

10716 rows × 10 columns

```
In [8]: print(len(trips))
```

10716

```
In [9]: # Use Seaborn `regplot()` function to plot the data and a linear regression model fit.
ax = sns.regplot(
    x="trip_distance", y="fare_amount",
    fit_reg=False, ci=None, truncate=True, data=trips)
ax.figure.set_size_inches(10, 8)
```



What's up with the streaks around 45 dollars and 50 dollars? Those are fixed-amount rides from JFK and La Guardia airports into anywhere in Manhattan, i.e. to be expected. Let's list the data to make sure the values look reasonable.

Let's also examine whether the toll amount is captured in the total amount.

```
In [10]: tollrides = trips[trips["tolls_amount"] > 0]
tollrides[tollrides["pickup_datetime"] == "2012-02-27 09:19:10 UTC"]
```

```
Out[10]:
```

	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	trip_distance	tolls_amount	fare_amount	total_amount
16	2012-02-27 09:19:10 UTC	-73.874431	40.774011	-73.983967	40.744082	1	11.6	4.8	27.7	

```
In [11]: notollrides = trips[trips["tolls_amount"] == 0]
notollrides[notollrides["pickup_datetime"] == "2012-02-27 09:19:10 UTC"]
```

```
Out[11]:
```

	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	trip_distance	tolls_amount	fare_amount	total_amount
52	2012-02-27 09:19:10 UTC	-73.972311	40.753067	-73.957389	40.817824	1	5.6	0.0	16.9	
7799	2012-02-27 09:19:10 UTC	-73.987582	40.725468	-74.016628	40.715534	1	2.8	0.0	12.1	
10537	2012-02-27 09:19:10 UTC	-74.015483	40.715279	-73.998045	40.756273	1	3.3	0.0	10.9	

Looking at a few samples above, it should be clear that the total amount reflects fare amount, toll and tip somewhat arbitrarily -- this is because when customers pay cash, the tip is not known. So, we'll use the sum of fare_amount + tolls_amount as what needs to be predicted. Tips are discretionary and do not have to be included in our fare estimation tool.

Let's also look at the distribution of values within the columns.

```
In [12]: # Print the distribution of values within the columns using `describe()`
trips.describe()
```

```
Out[12]:
```

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	trip_distance	tolls_amount	fare_amount	total_amount
count	10716.000000	10716.000000	10716.000000	10716.000000	10716.000000	10716.000000	10716.000000	10716.000000	10716.000000
mean	-72.602192	40.002372	-72.594838	40.002052	1.650056	2.856395	0.226428	11.109446	13.217078
std	9.982373	5.474670	10.004324	5.474648	1.283577	3.322024	1.135934	9.137710	10.953156
min	-74.258183	0.000000	-74.260472	0.000000	0.000000	0.010000	0.000000	2.500000	2.500000
25%	-73.992153	40.735936	-73.991566	40.734310	1.000000	1.040000	0.000000	6.000000	7.300000
50%	-73.981851	40.753264	-73.980373	40.752956	1.000000	1.770000	0.000000	8.500000	10.000000
75%	-73.967400	40.767340	-73.964142	40.767510	2.000000	3.160000	0.000000	12.500000	14.600000
max	0.000000	41.366138	0.000000	41.366138	6.000000	42.800000	16.000000	179.000000	179.000000

Copyright 2020 Google Inc. Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at <http://www.apache.org/licenses/LICENSE-2.0> Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.